

# homework\_1

October 10, 2021

## 1 Tasks — Regression (week 1):

### 1.0.1 Task 2:

```
[134]: import pandas as pd
import numpy as np

reviews = pd.read_json('https://cseweb.ucsd.edu/classes/fa21/cse258-b/data/
↳fantasy_10000.json.gz', lines=True)
```

```
[135]: ratings = [review for review in reviews["rating"]]
lengths = [len(review) for review in reviews["review_text"]]

y = np.matrix(ratings).T
X = np.matrix([[1,length] for length in lengths])

theta = np.linalg.inv(X.T*X)*X.T*y
print("Theta([Theta_0; Theta_1]): \n", theta)
```

```
Theta([Theta_0; Theta_1]):
[[3.68568136e+00]
 [6.87371675e-05]]
```

```
[136]: sse = sum([x**2 for x in (y - X*theta)]) #Summing up all square errors
mse = sse / len(y) #Dividing by length to get mean square error
print("MSE = ", mse)
```

```
MSE = [[1.55220866]]
```

### 1.0.2 Task 3:

```
[137]: import dateutil.parser
dates = [dateutil.parser.parse(review_time) for review_time in
↳reviews["date_added"]]
```

③ To use the piecewise function implementation of one-hot encoding the weekdays, we need a  $1 \times 6$  vector:

$[0 \ 0 \ 0 \ 0 \ 0 \ 0]$  it Monday

$[1 \ 0 \ 0 \ 0 \ 0 \ 0]$  it Tuesday

$[0 \ 1 \ 0 \ 0 \ 0 \ 0]$  it Wednesday

$\vdots$

$[0 \ 0 \ 0 \ 0 \ 0 \ 1]$  it Sunday

To one-hot encode the years in the same general manner is a bit worse, since there in theory might be an infinite amount of years. However, since I know the years in this dataset range from 2006 to 2017, I took the liberty of limiting the encoding to these years. We then need a  $1 \times 11$  vector:

$[0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0]$  it 2006

$[1 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0]$  it 2007

$[0 \ 1 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0]$  it 2008

$\vdots$

$[0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 1]$  it 2017

The feature vectors for the first two samples then becomes as follows:

$$\begin{aligned}
 x_1 &= [1 \quad \underbrace{2086}_{\text{length}} \quad \underbrace{0 \ 0 \ 0 \ 0 \ 0 \ 1}_{\text{day(sunday)}} \quad \underbrace{0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 1}_{\text{year(2017)}}] \\
 x_2 &= [1 \quad \underbrace{1521}_{\text{length}} \quad \underbrace{0 \ 1 \ 0 \ 0 \ 0 \ 0}_{\text{day(wednesday)}} \quad \underbrace{0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 1 \ 0 \ 0 \ 0}_{\text{year(2014)}}]
 \end{aligned}$$

### 1.0.3 Task 4:

```
[138]: #Using values directly as features
def feature(index):
    feature = [1]
    feature.append(lengths[index])
```

```

        feature.append(dates[index].weekday())
        feature.append(dates[index].year-2006)
    return feature

y = np.matrix(ratings).T
X = np.matrix([feature(index) for index in range(0, len(lengths))])

theta = np.linalg.inv(X.T*X)*X.T*y
print("Theta_1: \n", theta[0],
      "\nTheta_2: \n", theta[1],
      "\nTheta_3: \n", theta[2],
      "\nTheta_4: \n", theta[3])

sse = sum([x**2 for x in (y - X*theta)]) #Summing up all square errors
mse = sse / len(y) #Dividing by length to get mean square error
print("MSE = ", mse)

```

```

Theta_1:
[[3.29014782]]
Theta_2:
[[5.50923292e-05]]
Theta_3:
[[0.00875072]]
Theta_4:
[[0.05235923]]
MSE = [[1.53677405]]

```

```

[139]: #Using one-hot encoding
days = [date.weekday() for date in dates]
years = [date.year for date in dates]

def oneHotFeature(index):
    feature = [1]

    length = lengths[index]
    feature.append(length)

    day = dates[index].weekday()
    dayEncoded = [0]*6
    if day != 0:
        dayEncoded[day-1] = 1
    for elem in dayEncoded:
        feature.append(elem)

    year = dates[index].year
    yearEncoded = [0]*11
    if year != 2006:

```

```

        yearEncoded[year-2007] = 1
    for elem in yearEncoded:
        feature.append(elem)

    return feature

y = np.matrix(ratings).T
X = np.matrix([oneHotFeature(index) for index in range(0, len(lengths))])

theta = np.linalg.inv(X.T*X)*X.T*y
print("Theta_1: \n", theta[0],
      "\nTheta_2: \n", theta[1],
      "\nTheta_3: \n", theta[2:8],
      "\nTheta_4: \n", theta[8:])

sse = sum([x**2 for x in (y - X*theta)]) #Summing up all square errors
mse = sse / len(y) #Dividing by length to get mean square error
print("MSE = ", mse)

```

```

Theta_1:
[[4.87171479]]
Theta_2:
[[5.15709386e-05]]
Theta_3:
[[0.04890034]
 [0.1457098 ]
 [0.1066464 ]
 [0.12616832]
 [0.03834177]
 [0.1028469 ]]
Theta_4:
[[-1.58244783]
 [-1.70447417]
 [-1.68316056]
 [-1.67023905]
 [-1.62877001]
 [-1.19956705]
 [-1.10444816]
 [-1.09162361]
 [-1.20861354]
 [-1.23647487]
 [-1.23331225]]
MSE = [[1.51235787]]

```

#### 1.0.4 Task 5:

```
[147]: # distributing "random" indices to each set
training_indices = [i for i in range(0, len(lengths), 2)]
test_indices = [i for i in range(1, len(lengths), 2)]

training_ratings = [ratings[i] for i in training_indices]
test_ratings = [ratings[i] for i in test_indices]
```

```
[148]: #Using values directly as features
training_y = np.matrix(training_ratings).T
test_y = np.matrix(test_ratings).T
training_X = np.matrix([feature(index) for index in training_indices])
trained_theta = np.linalg.inv(training_X.T*training_X)*training_X.T*training_y

sse = sum([x**2 for x in (training_y - training_X*trained_theta)]) #Summing up
    ↪ all square errors
mse = sse / len(test_y) #Dividing by length to get mean square error
print("MSE(training data) = ", mse)
sse = sum([x**2 for x in (test_y - training_X*trained_theta)]) #Summing up all
    ↪ square errors
mse = sse / len(test_y) #Dividing by length to get mean square error
print("MSE(test data) = ", mse)
```

```
MSE(training data) = [[1.54991753]]
MSE(test data) = [[1.52654798]]
```

```
[149]: #Using one-hot encoding
training_y = np.matrix(training_ratings).T
test_y = np.matrix(test_ratings).T
training_X = np.matrix([oneHotFeature(index) for index in training_indices])
trained_theta = np.linalg.inv(training_X.T*training_X)*training_X.T*training_y

sse = sum([x**2 for x in (training_y - training_X*trained_theta)]) #Summing up
    ↪ all square errors
mse = sse / len(test_y) #Dividing by length to get mean square error
print("MSE(training data) = ", mse)
sse = sum([x**2 for x in (test_y - training_X*trained_theta)]) #Summing up all
    ↪ square errors
mse = sse / len(test_y) #Dividing by length to get mean square error
print("MSE(test data) = ", mse)
```

```
MSE(training data) = [[1.52069864]]
MSE(test data) = [[1.51164629]]
```

### 1.0.5 Task 6:

$$\textcircled{6} \quad \text{MAE} = \sum_{i=1}^{10000} |y_i - \theta_0|$$

$$\frac{\partial}{\partial \theta_0} \text{MAE} = - \underbrace{\sum_{i=1}^{10000} \delta(y_i > \theta_0)}_{y_i > \theta_0} + \underbrace{\sum_{i=1}^{10000} \delta(y_i < \theta_0)}_{y_i < \theta_0} = 0$$

If  $y_i > \theta_0$  and  $y_i < \theta_0$  happens exactly the same amount of times, the two terms above will cancel each other out, and  $\frac{\partial}{\partial \theta_0} \text{MAE}$  will be 0.

This has to be a global minima, since it does not make sense that increasing or decreasing  $\theta_0$  to the extreme will lead to a smaller MAE, which it would have done if it was a global maxima or a saddle point.

The only way to make the two terms cancel each other out, i.e. that  $y_i > \theta_0$  and  $y_i < \theta_0$  occurs the same amount of times, is if  $\theta_0$  has the value of the median of  $y$ .

## 2 Tasks — Classification (week 2):

### 2.0.1 Task 7:

```
[150]: import pandas as pd
import numpy as np
import sklearn
from sklearn import linear_model
from urllib.request import urlopen

def readFromURL(URL):
    for l in urlopen(URL):
        yield eval(l)

data = list(readFromURL("https://cseweb.ucsd.edu/classes/fa21/cse258-b/data/
→beer_50000.json"))
reviews = pd.DataFrame(data)
```

```
[152]: lengths = [len(review) for review in reviews["review/text"]]
X = np.matrix([[1,length] for length in lengths])
y = [rating >= 4 for rating in reviews["review/overall"]]

model = sklearn.linear_model.LogisticRegression(class_weight='balanced')
model.fit(X,y)
predictions = model.predict(X)

truePos = sum([(predictions[i] and y[i]) for i in range(0, len(y))])
fakePos = sum([(predictions[i] and not y[i]) for i in range(0, len(y))])
trueNeg = sum([(not predictions[i] and not y[i]) for i in range(0, len(y))])
fakeNeg = sum([(not predictions[i] and y[i]) for i in range(0, len(y))])

FPR = fakePos/(fakePos+trueNeg)
FNR = fakeNeg/(fakeNeg+truePos)

BER = 0.5*(FPR + FNR)

print("True Positive = ", truePos)
print("Fake Positive = ", fakePos)
print("True Negative = ", trueNeg)
print("Fake Negative = ", fakeNeg)
print("Balanced Error Rates = ", BER)
```

```
True Positive = 14201
Fake Positive = 5885
True Negative = 10503
Fake Negative = 19411
Balanced Error Rates = 0.46830315259572763
```




### 2.0.2 Task 8:

```
[154]: import matplotlib.pyplot as plt

confidences = model.decision_function(X)
sortedByConfidence = list(zip(confidences, y))
sortedByConfidence.sort(reverse=True)

K = [K for K in range(1, 10000+1, 10)]
precisionsK = []
for k in K:
    retrievedLabels = [x[1] for x in sortedByConfidence[:k]]
    precisionK = sum(retrievedLabels)/len(retrievedLabels)
    precisionsK.append(precisionK)

plt.plot(K, precisionsK)
plt.xlabel("K")
plt.ylabel("Precision")
plt.show()
print(f"precision@K for k=1: {precisionsK[0]}, k=100: {precisionsK[100//10-1]}, k=10000: {precisionsK[10000//10-1]}")
```



precision@K for k=1: 1.0, k=100: 0.7692307692307693, k=10000: 0.7152437193474127

### 2.0.3 Task 9:

```
[155]: import matplotlib.pyplot as plt


confidences = model.decision_function(X)
sortedByConfidence = list(zip(abs(confidences), y, predictions))
sortedByConfidence.sort(reverse=True)

K = [K for K in range(1, 10000+1, 10)]
precisionsK = []
for k in K:
    retrievedLabels = [[x[1], x[2]] for x in sortedByConfidence[:k]]
    precisionK = 0
    for elem in retrievedLabels:
        precisionK += (int)(elem[0]==elem[1])
    precisionK = precisionK/k
    precisionsK.append(precisionK)

plt.plot(K, precisionsK)
plt.xlabel("K")
plt.ylabel("Precision")
plt.show()

print(f"precision@K for k=1: {precisionsK[0]}, k=100: {precisionsK[100//10-1]},  

    ↪k=10000: {precisionsK[10000//10-1]}")
```



precision@K for k=1: 1.0, k=100: 0.7692307692307693, k=10000: 0.6188569712741467